

CS111 - Project 3A: File System Dump

INTRODUCTION:

Project 3 is expected to involve much more programming than any of the other projects. You will design and implement programs to analyze file systems and diagnose corruption. In part A, you will write a program to read the image of a file system, analyze it, and summarize its contents in several csv files. These csv files will be used in part B, where we will analyze the consistency of file system data structures and diagnose problems.

Part A can be broken into two major steps:

1. Mount the provided image file on your own Linux and explore it with *debugfs(8)*.
2. Write a program to analyze the image file and output a summary to six csv files (describing the super block, cylinder groups, free-lists, inodes, indirect blocks, and directories).

RELATION TO READING AND LECTURES:

This project more deeply explores the file and directory concepts described in Arpaci chapter 39.

This project is based on the same EXT2 file system that is discussed in sections 40.2-40.5. Part B of this project goes much deeper than the introductory discussion of integrity presented in sections 42.1-2.

PROJECT OBJECTIVES:

- reinforce the basic file system concepts of directory objects, file objects, and free space.
- reinforce the implementation descriptions provided in the text and lecture.
- gain experience with the examining, interpreting, and processing information in binary data structures.
- gain practical experience with complex data structures in general, and on-disk data formats in particular.
- (in 3B) reinforce the notions of consistency/integrity provided in the text and lecture.

DELIVERABLES:

A single tarball (.tar.gz) containing:

- a single C source module that compiles cleanly (with no errors or warnings).
- a Makefile to build the program and the tarball.
- a README file describing each of the included files, your UID, and any other information about your submission that you would like to bring to our attention (e.g. limitations, features, testing methodology, etc.).

PROJECT DESCRIPTION:

Historically, file systems were almost always been implemented as part of the operating system, running in kernel mode. Kernel code is expensive to develop, difficult to test, and prone to catastrophic failures. Within the past 15 years or so, new developments have made it possible to implement file systems in user mode, improving maintainability ... and in some cases delivering even better performance than could be achieved with kernel code. All of this project will be done as user-mode software.

To ensure data security and integrity, disks are generally protected from access by ordinary applications. Linux supports the creation, mounting, checking, and debugging of file systems stored in ordinary files. In this project, we will provide EXT2 file system images in ordinary files. Because they

are in ordinary files (rather than protected disks) you can access/operate on those file system images with ordinary user mode code.

To complete this assignment, you will need to learn a few things:

- *debugfs(8)* (<http://man7.org/linux/man-pages/man8/debugfs.8.html>)
- *pread(2)*
- csv format
- EXT2 file system (<http://www.nongnu.org/ext2-doc/ext2.html>)

Step 1 - study a provided file system image

Your program will be tested on multiple (broken) file systems. Here we provide one as an example. SEASnet servers do not support commands like *sudo(8)*, *mount(8)*, or *debugfs(8)*, so to play with our provided file system image, you will have to install a Linux distribution (if you do not have one) on your own computer. For simplicity and convenience, you may choose to install Linux inside a virtual machine, for example VirtualBox (it's free!).

Then download [this image file](#), and mount it onto your **own** Linux, with the following commands:

```
mkdir fs
sudo mount -o loop disk-image fs
sudo chown -R $USER fs          ... where $USER is your user-name
```

Note: this file is 200 Mbytes in size.

Now, you can navigate the file system, just like an ordinary directory, with commands like *ls(1)*, *cat(1)*, and *cd(1)*. After you are done with it, you can unmount with the following command:

```
sudo umount fs
```

Before you start writing your C program to interpret the “disk-image” file, you can explore it further using *debugfs(8)* (in your own Linux). You may find many useful commands in its man page. Some particularly helpful ones are: *stats*, *stat*, *bd*, *testi*, and *testb*. If you have problems interpreting parts of the file system, you can use the *debugfs* program to help you understand the correct contents.

To ensure you are correctly interpreting the file system image, we have included many unusual things, which would not likely be caught by an incorrect implementation:

- sparse files
- large files
- allocated data blocks full of zeroes
- unallocated blocks containing valid data
- files with data beyond their length
- files with long names
- files with syntactically strange names
- directories that span multiple blocks and have obsolete entries for deleted files.

Step 2 - write a program to summarize the file system's contents

In this step, you will write a program called **lab3a** that:

- Reads (only) one file system image according to the provided file name. For example, we may run your program with the above file system image using the following command: `./lab3a disk-image`
- Analyzes the provided file system image and outputs six csv files to the current directory. The contents of these csv files are described below. Your program **must** output these files with **exactly the same formats** as shown below. We will use sort and diff to compare your csv files with ours, so a different format will make your program fail the test and the content in the output csv file will be treated as error.

Please note that, although you cannot mount the provided image file on departmental servers, your lab3a program should be able to run on departmental servers, just like all previous assignments.

There are to be six csv-format files, each summarizing a different part of the file system. Remember, you can always check your program's output against *debugfs*'s output. All the information required for the summary can be manually found and checked by using debugfs.

The file system images your program examines will, in many cases be damaged (as might result from missing or mis-directed writes). This means that you will have to do some *Defensive Programming* in your analysis program: validating parameters and pointers before you use them. In addition to correctly analyzing the file system and group parameters, inodes, block pointers, and directory entries, your program will also be expected to report (to stderr) any invalid parameters or addresses.

1. super block

A basic set of file system parameters.

fields(9)	format
magic number	hex
total number of inodes	dec
total number of blocks	dec
block size	dec
fragment size	dec
blocks per group	dec
inodes per group	dec
fragments per group	dec
first data block	dec

Sample correct csv output

```
Ef53,51200,204800,1024,1024,8192,2048,8192,1
```

Sanity checking

- Magic number must be correct
- Block size must be reasonable (e.g. power of two between 512-64K)
- Total blocks and first data block must be consistent with the file size
- Blocks per group must evenly divide into total blocks
- Inodes per group must evenly divide into total inodes

Sample error messages (to stderr):

```
Superblock - invalid magic: 0xdead
Superblock - invalid block size: 666
Superblock - invalid block count 200000 > image size 50000
Superblock - invalid first block 100000 > image size 50000
Superblock - 20000 blocks, 1050 blocks/group
Superblock - 1000 Inodes, 66 Inodes/group
```

It is important to validate superblock parameters, because they are used to determine the addresses of inodes and to validate block pointers. If a file system fails any of these tests, log an error and exit without any further processing.

1. group descriptor

Information about each cylinder group

fields(7)	format
number of contained blocks	dec
number of free blocks	dec
number of free inodes	dec
number of directories	dec
(free) inode bitmap block	hex
(free) block bitmap block	hex
inode table (start) block	hex

8192,7913,2024,2,4,3,5

- Number of contained blocks must be consistent with superblock

Sample error messages (to stderr):

Group 7: blocks 100000-150000, free Inode map starts at 165000

Group 7: blocks 100000-150000, Inode table starts at 165000

1. free bitmap entry:

One line of output for each free inode or block in the bitmap. No output should be produced for allocated inodes or blocks.

block number of the map	hex
-------------------------	-----

Sample output (just first five entries)

3,276

3,278

3,279

- none ... since there are no block numbers within the bit-map, there is no trivial way to tell if the map has been corrupted.

Key information for each allocated inode.

fields(11+15) format

```
file type      char[1]
```

group	dec
-------	-----

creation time	hex
---------------	-----

access time	hex
1	00000000
2	00000000
3	00000000
4	00000000
5	00000000
6	00000000
7	00000000
8	00000000
9	00000000
10	00000000
11	00000000
12	00000000
13	00000000
14	00000000
15	00000000
16	00000000
17	00000000
18	00000000
19	00000000
20	00000000
21	00000000
22	00000000
23	00000000
24	00000000
25	00000000
26	00000000
27	00000000
28	00000000
29	00000000
30	00000000
31	00000000
32	00000000
33	00000000
34	00000000
35	00000000
36	00000000
37	00000000
38	00000000
39	00000000
40	00000000
41	00000000
42	00000000
43	00000000
44	00000000
45	00000000
46	00000000
47	00000000
48	00000000
49	00000000
50	00000000
51	00000000
52	00000000
53	00000000
54	00000000
55	00000000
56	00000000
57	00000000
58	00000000
59	00000000
60	00000000
61	00000000
62	00000000
63	00000000
64	00000000
65	00000000
66	00000000
67	00000000
68	00000000
69	00000000
70	00000000
71	00000000
72	00000000
73	00000000
74	00000000
75	00000000
76	00000000
77	00000000
78	00000000
79	00000000
80	00000000
81	00000000
82	00000000
83	00000000
84	00000000
85	00000000
86	00000000
87	00000000
88	00000000
89	00000000
90	00000000
91	00000000
92	00000000
93	00000000
94	00000000
95	00000000
96	00000000
97	00000000
98	00000000
99	00000000
100	00000000
101	00000000
102	00000000
103	00000000
104	00000000
105	00000000
106	00000000
107	00000000
108	00000000
109	00000000
110	00000000
111	00000000
112	00000000
113	00000000
114	00000000
115	00000000
116	00000000
117	00000000
118	00000000
119	00000000
120	00000000
121	00000000
122	00000000
123	00000000
124	00000000
125	00000000
126	00000000
127	00000000
128	00000000
129	00000000
130	00000000
131	00000000
132	00000000
133	00000000
134	00000000
135	00000000
136	00000000
137	00000000
138	00000000
139	00000000
140	00000000
141	00000000
142	00000000
143	00000000
144	00000000
145	00000000
146	00000000
147	00000000
148	00000000
149	00000000
150	00000000
151	00000000
152	00000000
153	00000000
154	00000000
155	00000000
156	00000000
157	00000000
158	00000000
159	00000000
160	00000000
161	00000000
162	00000000
163	00000000
164	00000000
165	00000000
166	00000000
167	00000000

number of blocks dec

```
block pointers * 15    hex
```

1,?,0,0,0,0,58201197,58201197,58201197,0,0,0,0,0,0,0,0,0,0,0,0,0,0

2,d,40755,1000,0,5,58201213,58201213,582011bd,3072,3,105,a602,a107,0,0,0,0,0,0,0,0,0,
0

Sanity checking

- It is important to validate all block pointers before we use them:

- Sample error message (to stderr):

If any block address falls outside the reasonable range, include it in the output, log an error message. But if an out-of-range block is an indirect block or contains directory entries, do not attempt to further interpret its contents.

The valid/allocated entries in each directory.

The term *parent inode number*, below, refers to the inode number of the directory that contains all of the entries that are being listed.

fields(6)	format
-----------	--------

name string[2]

```
2,2,20,10,11,"lost+found"
```

- file entry inode number should be within the super-block specified range (number of inodes)

1. indirect block entry

These are all the **non-zero** block pointers in an indirect block. The blocks that contain indirect block pointers are included.

fields(3)	format
block number of the containing block	hex
entry number within that block	dec
block pointer value	hex

Sample output (first three entries)

```
1e10f,0,1e110
1e10f,1,1e111
1e10f,2,1e112
```

Sanity checking

- all block numbers should be legal (within super-block specified range)

It is important to validate all block numbers in indirect blocks for the same reason we had to validate block numbers in inodes: because we will follow those pointers to further interpret the file system.

Sample error message to stderr:

```
Indirect block 3c4 - invalid entry[103] = 304c8df
```

If any block number is out of range log an error message. If a out-of-range block is an indirect block, do not attempt to dump its contents.

Output Files

For each different kind of entry, output its summary in a separate file. The names of your files should be:

- super.csv ...for super block;
- group.csv ...for group descriptors;
- bitmap.csv ...for free inodes and blocks;
- inode.csv ...for inodes;
- directory.csv ...for directory entries; and
- indirect.csv ...for indirect block entries.

In these csv files, each line represents one entry, and should be ended with a single new-line character ('\n', 0x0a). The field orders in those csv files should be the same as listed above, and your program should **use and only use** a comma to separate each field.

For your convenience, here are the exact six csv files generated by our solution: [super.csv](#), [group.csv](#), [bitmap.csv](#), [inode.csv](#), [directory.csv](#), and [indirect.csv](#).

SUBMISSION:

Please carefully check the Rubric section (especially the **Code review** subsection) and **all (latest) footnotes** before submitting your code.

Your tarball should have a name of the form `lab3a-studentID.tar.gz` and should be submitted via CCLE.

We will test your work on a SEASnet GNU/Linux server running RHEL 7 (this is on Inxsrv09). You would be well advised to test your submission on that platform before submitting it.

RUBRIC:

Value Feature

Packaging and build (10%)

- 3% untars expected contents
- 3% clean build w/default action (no warnings)
- 2% Makefile has clean and dist targets
- 2% reasonableness of README contents

Code review (15%)

You can check header files like `ext2_fs.h` to have an idea about the variable types. HOWEVER, you shall write up your own code for the calculations like block/inode (number), indirect block pointer, various different kinds of offset, and so on. So in sum, do all the calculations with your own code, do not use macros/functions provided in header files like `ext2fs.h`.

5% overall readability and reasonableness, correct program name and csv filenames

10% image file investigation correct (use `pread()` and write your own code for calculations)

Results (75%)

$\frac{2}{3}$ points shall go to outputting the required fields correctly, and $\frac{1}{3}$ points go to correct lines (penalty will be given for extra/missing lines).

10% `super.csv`

10% `group.csv`

10% `bitmap.csv`

15% `inode.csv`

15% `directory.csv`

15% `indirect.csv`

[1] We only recognize a (small) subset of the file types: 'f' for regular file, 'd' for directory, and 's' for symbolic links. For other types, use '?'.

[2] Double-quoted. Don't worry about how to print out non-graphical characters or embedded quotes.